

Homework 3

Matt Dickenson
CS 527
Fall 2014

Problem 1

(a) kmeans.m

```
function [M, R] = kmeans(Z, M)
    % initialize f(S) values for iteration
    f_p = intmax;
    f_c = f_p - 1;

    while f_p > f_c + sqrt(eps)
        R = closest_cluster(Z, M);
        M = (Z * transpose(R)) ./ (transpose(sum(R, 2) * ones(1, 2)));
        f_p = f_c;
        f_c = sum(sum(dists(Z, M) .* R, 1), 2);
    end
end

function [assignments] = closest_cluster(Z, mu)
    % partition the points in Z by their closest mean
    K = size(mu, 2);
    all_dists = dists(Z, mu);
    min_dists = ones(K, 1) * min(all_dists, [], 1);
    assignments = all_dists == min_dists;
end

function [all_dists] = dists(Z, mu)
    % compute distance of each point in Z from each centroid
    K = size(mu, 2);
    I = size(Z, 2);
    all_dists = zeros(K, I);
    for k=1:K
        means = mu(:, k) * ones(1, I);
        dists = sqrt(sum((Z - means).^2, 1));
        all_dists(k, :) = dists;
    end
end
```

Output:

M =

0.6816	0.8271	-1.0890
-0.6389	1.1812	-0.6784




Figure 1: Result of K-Means with blobs data for 1(a)

(b) Output:

M =

-1.0371	1.0207
-0.0176	0.0173




Figure 2: Result of K-Means with cigars data for 1(b)

(c) Output:

$M =$

-1.0402	1.0175
-0.0199	0.0195




Figure 3: Result of K-Means with cigars data for 1(c)

(d) The shortcomings are not that the K-means algorithm computes a local minimum. The problem is that both dimensions are treated equally for the purposes of computing the error rate despite the fact that they may have very different variances.

The problem is that the data points vary much more along the x-dimension than the y dimension. Although they appear to be well-described by two clusters with the same x range, one above the other, the high variance in the x dimension would force two such centroids to still have a high error measure $f(S)$.

(e) I constructed four clusters whose points form approximately the outline of a unit square. The four initial centroids m all lie at the origin. The points




Figure 4: Result of K-Means with pathological input

and centroids are displayed in Figure 4 below. Note that the final centroids and initial centroids are exactly equal.

The true centroids are $\mu_1^* = (0, 1), \mu_2^* = (1, 0), \mu_3^* = (0, -1), \mu_4^* = (-1, 0)$. The initial (and final) centroids are $m_1 = m_2 = m_3 = m_4 = (0, 0)$.

The points in each cluster are

$$\begin{aligned} S_1 &= \{(\epsilon, 1), (1 + \epsilon, 0), (-\epsilon, -1), (-1 - \epsilon, 0)\} \\ S_2 &= \{(-\epsilon, 1), (1 - \epsilon, 0), (\epsilon, -1), (-1 + \epsilon, 0)\} \\ S_3 &= \{(0, 1 + \epsilon), (1, \epsilon), (0, -1 - \epsilon), (1, -\epsilon)\} \\ S_4 &= \{(0, 1 - \epsilon), (1, -\epsilon), (0, -1 + \epsilon), (-1, \epsilon)\}, \end{aligned}$$

although any random, symmetric perturbation about the true centroids would work in expectation.

The algorithm makes no progress because the initial centroids m are each symmetrically located between all four clusters and all clusters are equally sized. Moving any centroid nearer to one cluster would make it closer to a few points but farther away from all other points, cancelling out the advantage of relocating the centroid.

(f) The previous answer shows that the success K -means algorithm depends upon the initial centroids. K -means also depends on knowing the true number of clusters. It can also fail when clusters are symmetric and of equal size.

Problem 2

(a) EM.m

```
function [lambda, M, Sigma, R] = EM(Z, M, sigma2)
    [d, I] = size(Z);
    K = size(M, 2);

    % Initialize uninformative prior
    lambda = repmat(1/K, K, 1);
    Sigma = zeros(d, d, K);
    for k=1:K
        Sigma(:, :, k) = sigma2(k)*eye(d);
    end
    R = zeros(K, I);

    % initialize f(theta) values for iteration
```

```

f_p = intmax;
f_c = f_p - 1;

while f_p > f_c + sqrt(eps)
    % E-step
    for k=1:K
        for i=1:I
            R(k, i) = lambda(k) * mymvn(Z(:, i), M(:, k), Sigma(:, :, k));
        end
    end
    R = R ./ (ones(K, 1) * sum(R));

    % M-step
    for k=1:K
        lambda(k) = mean(R(k, :));
        r_d = ones(d, 1) * R(k, :);
        M(:, k) = sum( r_d .* Z, 2 ) / sum(R(k, :));
        diff = Z - (M(:, k) * ones(1, I));
        Sigma(:, :, k) = ((r_d .* diff) * diff') / sum(R(k, :));
    end

    f_p = f_c;
    f_c = -sum(log(sum(R, 2)));
end
end

function [p] = mymvm(x, mu, sigma)
    d = size(x, 1);
    if size(sigma) ~= [d, d]
        error('Sigma must be of dimension dxd')
    end
    normalization = 1/sqrt((2*pi)^d * det(sigma));
    exponent = (x-mu)' * inv(sigma) * (x-mu);
    p = normalization * exp(-0.5*exponent);
end

```

Output:

```

lambda =
    0.3031
    0.3638
    0.3331

M =
    0.2733    0.8352   -1.1607

```

```

-0.7696    1.1658   -0.5729

Sigma(:,:,1) =
    0.5661    0.1155
    0.1155    0.1176

Sigma(:,:,2) =
    0.1578    0.0080
    0.0080    0.1610

Sigma(:,:,3) =
    0.2156    0.0474
    0.0474    0.2841

```




Figure 5: Result of EM with blobs data for 2(a)

(b) For the cigars data, the E-M result is:

```
lambda =
    0.5000
    0.5000

M =

   -0.0333    0.0333
   -0.5701    0.5701

Sigma(:,:,1) =
    1.6468   -0.0001
   -0.0001    0.0175

Sigma(:,:,2) =
    1.6652    0.0007
    0.0007    0.0162
```




Figure 6: Result of EM with cigars data for 2(b)

(c) Yes, the result of E-M using the cigars data is much more satisfactory than using K -means. Unlike with K -means, the centroids are now each centered on one of the two clusters in the data. There is also little uncertainty about which cluster each point belongs to.

(d) For E-M on the blobs data with $k=2$ the result is:

With $k = 4$ the results are:

(e) For the bananas data, the E-M result is:

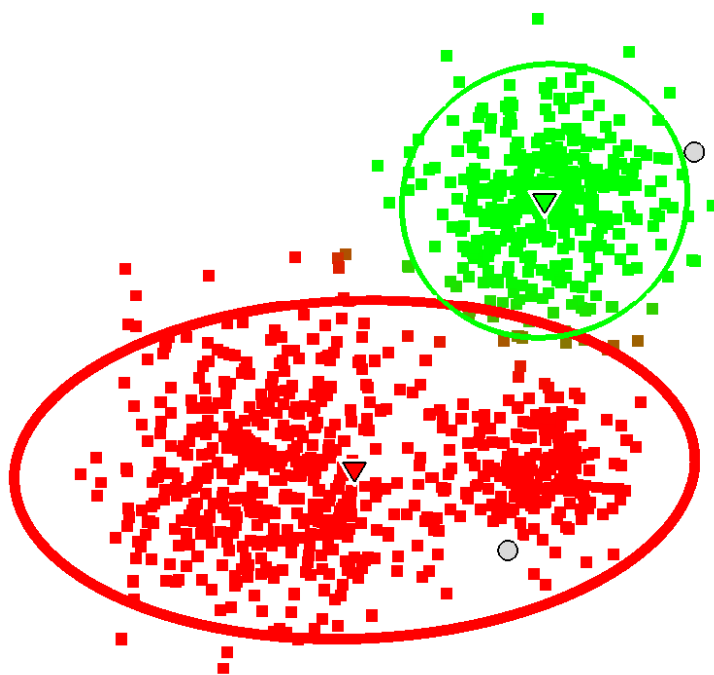


Figure 7: Result of EM with blobs data and $k=2$




Figure 8: Result of EM with blobs data and $k=4$

M =

-1.0771	1.0769
-0.6489	0.6488




Figure 9: Result of EM with bananas data for 2(d)

(f) The result in 2(d) is unsatisfactory for two reasons. First, points at the top tip of the left banana and the bottom tip of the right banana are each misclassified as belonging to the other banana. This is a result of using a mixture of normal distributions, since the clusters are constrained to be ellipse-shaped.

Second, the ellipses defined by the mean and variance of each cluster exclude a number of points belonging to the cluster. This is a result of the E-M algorithm: this data could likely be well-described if $k = 6$, but it's difficult to know this a priori and k is fixed in this E-M implementation.

Problem 3

(a) `meanShift.m`:

```
function [z, zh] = meanShift(zstart, Z, h)
    [d, I] = size(Z);
    z_prime = zstart;
    zh = z_prime;
    terminate = false;
    while ~terminate
        z = z_prime;
        diff = z * ones(1, I) - Z;
        norms = sqrt(sum(diff.^2, 1));
        kernel = exp(-(norms/h).^2);
        z_prime = (Z * kernel') ./ sum(kernel, 2);
        zh = [zh'; z_prime'];
        terminate = sqrt(sum((z-z_prime).^2, 1)) <= h/1000;
    end
end
```

Running this code as instructed produced the following plot:

The first run with $h = 0.2$ took 36 iterations and the second run with $h = 2$ took 13 iterations.

(b) With $h = 2$ we have a larger neighborhood size, so instead of finding a local mode we find the global mode for the data given. Increasing h further results in approximately the same final point.

Problem 4

(a)

My implementation of the mean-shift clustering algorithm, with an improvement to speed up mean-shift, is shown below.

`meanShiftCluster.m`:

```
function [U, R] = meanShiftCluster(Z, h)
    [d, I] = size(Z);
    M = zeros(d, I);
    for i=1:I
        zstart = Z(:, i);
        idx = rangesearch(zstart', Z', h/100);
```

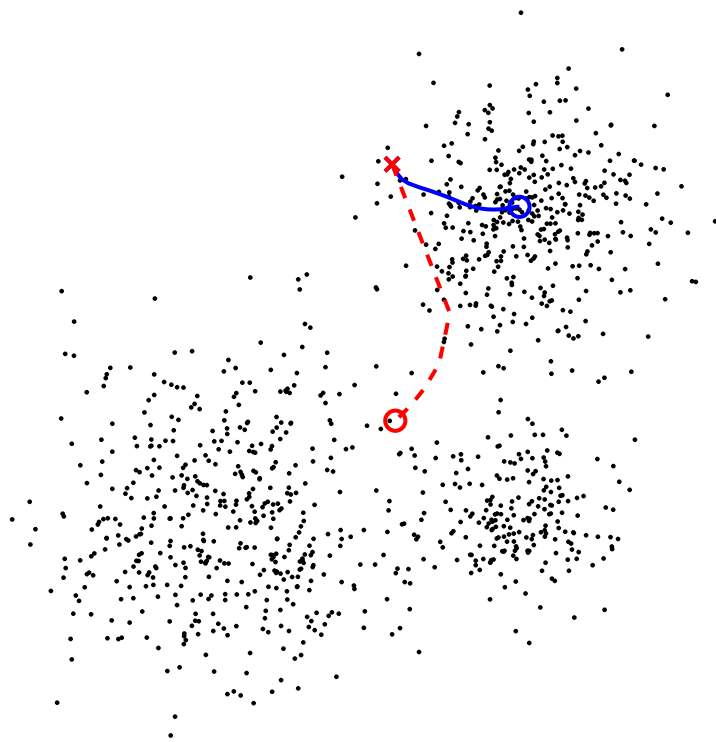


Figure 10: Result of Mean-Shift with blobs data

```

        keep = cellfun(@isempty, idx);
        tmpZ = Z(:, keep);
        [d, I] = size(tmpZ);
        z_prime = zstart;
        terminate = false;
        while ~terminate
            z = z_prime;
            diff = z * ones(1, I) - tmpZ;
            norms = sqrt(sum(diff.^2, 1));
            kernel = exp(-(norms/h).^2);
            z_prime = (tmpZ * kernel') ./ sum(kernel, 2);
            terminate = sqrt(sum((z-z_prime).^2, 1)) <= h/1000;
        end
        M(:, i) = z;
    end

U = nearUniqueCols(M, h/100);
K = size(U, 2);
R = false(K, I);
for k=1:K
    for i=1:I
        if sqrt(sum((M(:, i) - U(:, k)).^2)) < h/100
            R(k, i) = 1;
        else
            R(k, i) = 0;
        end
    end
end
end

function [u] = nearUniqueCols(mx, tau)
    keep = mx(:, 1);
    for j=2:size(mx, 2)
        col = mx(:, j);
        keep = unique(keep, col, tau);
    end
    u = keep;
end

function [u] = unique(mx, x, tau)
    for l=1:size(mx, 2)
        col = mx(:, l);
        dist = sqrt(sum((col-x).^2));
        if dist < tau
            u = mx;
            return
        end
    end
end

```

```

    end
  end
  u = [mx'; x']';
end

```

The running time of my code was 14.91 seconds (using the `timeit` function) on a 2-core 2010 Macbook Pro running OS X 10.9. The algorithm found two clusters, shown below.




Figure 11: Result of Mean-Shift Clustering with bananas data and $h=0.8$

(b) In this example, mean-shift clustering does not appear to do substantially better than EM. Again, the upper-right points of the left banana and the lower-

left points of the right banana are misclassified. One advantage of mean-shift clustering over EM is that we did not have to specify the number of clusters a priori. One disadvantage is that the mean-shift clustering results do not indicate uncertainty about classification as the EM results do.

Although we do not have to select k for mean-shift clustering, the selection of h is accompanied by trade-offs. A low bandwidth (small h) is likely to find a larger number of clusters centered at local modes. A larger bandwidth (greater h) will find fewer clusters, perhaps only one centered at the global mode. Thus, the selection of h presents similar challenges as the selection of k for K -means or EM.